



华南理工大学

South China University of Technology

The Experiment Report of *Machine Learning*

College Software College

Subject Software Engineering

Members _____

Student ID 201530661741

E-mail yimingzhao0@qq.com

Tutor Mingkui Tan

Date submitted 2017. .

1. Topic: Logistic regression, logistic classification and gradient decent method

2. Time: 2017.12.2 PM 2:00-5:00

3. Reporter: Yiming Zhao

4. Purposes:

- **Understand the differences between the gradient decent and stochastic gradient decent method.**
- **Understand the connection between the logistic regression and linear regression.**
- **Practicing the SVM on the bigger data set.**

5. Data sets and data analysis:

Using a9a data set from LIBSVM Data, including 32561 / 16281(testing) samples and 123 features in each sample.

6. Experimental steps:

- **Read the data set using sklearn's function**
- **Randomly initializing the parameters.**
- **Selecting a proper loss function and using NAG, Adam, Adadelta, RMSProp to optimize the loss function.**

7. Code:

```
from sklearn import datasets as ds
import numpy as np
```

```
from numpy import random
import matplotlib.pyplot as plt

def sigmoid(input_):
    return 1 / (1 + np.exp(-input_))

def train(x_train, y_train, x_test, y_test,
method, iters, test_errors):
    max_iterations = 100
    theta = random.rand(num_features + 1)
    num_test_samples, num_test_features =
x_test.shape

    if method == 'sgd':
        lr = 0.01

        for i in range(max_iterations):
            output = sigmoid(np.dot(x_train[i],
theta))
            error = output - y_train[i]
```

```

        theta = theta - lr * np.dot(x_train[i],
error)

        predict_error = 0
        for j in range(num_test_samples):
            predict_output =
sigmoid(np.dot(x_test[j], theta))
            predict_error -= y_test[j] *
np.log(predict_output) + (1 - y_test[j]) *
np.log(1 - predict_output)
            print(str(i) + '\t' +
str(predict_error / num_test_samples))

        iters.append(i)
        test_errors.append(predict_error /
num_test_samples)

    if method == 'nag':
        lr = 0.01
        miu = 0.9
        momentum = np.zeros(num_features + 1)

```

```
for i in range(max_iterations):  
    output = sigmoid(np.dot(x_train[i],  
theta - lr * miu * momentum))  
    error = output - y_train[i]  
    grad = np.dot(x_train[i], error)  
    momentum = momentum * lr + grad  
    theta = theta - lr * momentum  
  
    predict_error = 0  
    for j in range(num_test_samples):  
        predict_output =  
sigmoid(np.dot(x_test[j], theta))  
        predict_error -= y_test[j] *  
np.log(predict_output) + (1 - y_test[j]) *  
np.log(1 - predict_output)  
        print(str(i) + '\t' +  
str(predict_error / num_test_samples))  
  
    iters.append(i)  
    test_errors.append(predict_error /  
num_test_samples)
```

```

if method == 'rmsprop':
    lr = 0.1
    expectation = 1
    rho = 0.95
    delta = 10e-7

    for i in range(max_iterations):
        output = sigmoid(np.dot(x_train[i],
theta))

        error = output - y_train[i]
        grad = np.dot(x_train[i], error)
        norm = grad * grad
        expectation = rho * expectation + (1 -
rho) * norm

        theta = theta - lr * grad /
(np.sqrt(expectation) + delta)

        predict_error = 0
        for j in range(num_test_samples):
            predict_output =
sigmoid(np.dot(x_test[j], theta))
            predict_error -= y_test[j] *

```

```

np.log(predict_output) + (1 - y_test[j]) *
np.log(1 - predict_output)
        print(str(i) + '\t' +
str(predict_error / num_test_samples))
        iters.append(i)
        test_errors.append(predict_error /
num_test_samples)

if method == 'adam':
    delta = 10e-8
    rho1 = 0.9
    rho2 = 0.999
    lr = 0.1
    s = 0
    r = 0

    for i in range(max_iterations):
        output = sigmoid(np.dot(x_train[i],
theta))

        error = output - y_train[i]
        grad = np.dot(x_train[i], error)

```

```

        s = rho1 * s + (1 - rho1) * grad
        r = rho2 * r + (1 - rho2) * grad * grad
        s_hat = s / (1 - rho1)
        r_hat = r / (1 - rho2)
        delta_theta = (-lr * s_hat) /
(np.sqrt(r_hat) + delta)
        theta = theta + delta_theta

    predict_error = 0
    for j in range(num_test_samples):
        predict_output =
sigmoid(np.dot(x_test[j], theta))
        predict_error -= y_test[j] *
np.log(predict_output) + (1 - y_test[j]) *
np.log(1 - predict_output)
        print(str(i) + '\t' +
str(predict_error / num_test_samples))
        iters.append(i)
        test_errors.append(predict_error /
num_test_samples)

    if method == 'adadelata':

```



```

r = 0
e = 0
miu = 0.9
delta = 10e-7
lr = 10

for i in range(max_iterations):
    output = sigmoid(np.dot(x_train[i],
theta))

    error = output - y_train[i]
    grad = np.dot(x_train[i], error)

    r = miu * r + (1 - miu) * grad * grad
    delta_theta = (-lr * grad * np.sqrt(e
+ delta)) / (np.sqrt(r + delta))
    theta = theta + delta_theta
    e = miu * e + (1 - miu) * e * e

predict_error = 0
for j in range(num_test_samples):
    predict_output =

```

```

sigmoid(np.dot(x_test[j], theta))
        predict_error -= y_test[j] *
np.log(predict_output) + (1 - y_test[j]) *
np.log(1 - predict_output)
        print(str(i) + '\t' +
str(predict_error / num_test_samples))
        iters.append(i)
        test_errors.append(predict_error /
num_test_samples)

if __name__ == '__main__':
    x_train, y_train =
ds.load_svmlight_file('./data/a9a')
    x_test, y_test =
ds.load_svmlight_file('./data/a9a.t')

    num_samples, num_features = x_train.shape
    num_test_samples, num_test_features =
x_test.shape

    x_train = x_train.toarray()
    temp = np.ones(shape=[32561, 1],

```

```
dtype=np.float32)

    x_train = np.concatenate([x_train, temp],
axis=1)

    x_test = x_test.toarray()

    temp = np.zeros(shape=[16281, 1],
dtype=np.float32)

    temp1 = np.ones(shape=[16281, 1],
dtype=np.float32)

    x_test = np.concatenate([x_test, temp, temp1],
axis=1)


    for i in range(0, len(y_train)):
        if y_train[i] == -1:
            y_train[i] = 0
    for i in range(0, len(y_test)):
        if y_test[i] == -1:
            y_test[i] = 0


    methods = ['sgd', 'nag', 'rmsprop',
'adadelata', 'adam']

    for method in methods:
```

```
    iters = []
    test_errors = []
    train(x_train, y_train, x_test, y_test,
method, iters, test_errors)

    plt.plot(iters, test_errors,
label=method)

plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

```
from sklearn import datasets as ds
import numpy as np
from numpy import random
import matplotlib.pyplot as plt
import os

C = 1
feature_size = 123
bias = np.zeros(shape=[feature_size + 1, 1])
bias[len(bias)-1][0] = 1.
```

```
def compute_loss(x, y, w):
    # x.shape = [batch_size, feature_size + 1],
    y.shape = [batch_size, 1], w.shape =
    [feature_size + 1, 1]
    pred = np.matmul(x, w)
    hinge_loss = np.maximum(1 - y * pred, 0)
    loss = np.mean(hinge_loss ** 2) + C * np.sum((w
- bias) ** 2)
    return loss
```

```
def compute_gradient(x, y, w):
    # x.shape = [batch_size, feature_size + 1],
    y.shape = [batch_size, 1], w.shape =
    [feature_size + 1, 1]
    pred = np.matmul(x, w)
    hinge_loss = np.maximum(1 - y * pred, 0)
    hinge_loss_gradient =
    -np.matmul(x.transpose(), hinge_loss * y) / len(y)
    norm_gradient = 2 * C * (w - bias)
```

```

    gradient = hinge_loss_gradient +
norm_gradient

    return gradient # gradient.shape =
[feature_size + 1, 1], the same as w

global_list = {'NAG_momentum': 0,
'RMS_expectation': 1, 'ADAM_s': 0, 'ADAM_r': 0,
               'ADADELDA_r': 0, 'ADADELDA_e': 0}

def optimizer(method, parameter_list, x, y, w):
    global global_list
    if method == 'SGD':
        # sgd_para_list = {'learning_rate':0.01}
        lr = parameter_list['learning_rate']
        w -= lr * compute_gradient(x, y, w)
    if method == 'NAG':
        # nag_para_list = {'miu':0.9,
'learning_rate':0.01}
        lr = parameter_list['learning_rate']
        miu = parameter_list['miu']

```

```

        momentum = global_list['NAG_momentum']
        gradient = compute_gradient(x, y, w -
momentum * lr * miu)

        momentum = momentum * lr + gradient
        w -= lr * momentum
        global_list['NAG_momentum'] = momentum
    if method == 'RMSProp':
        # rmsprop_para_list = {'delta':10e-7,
'rho':0.95, 'learning_rate':0.1}

        lr = parameter_list['learning_rate']
        expectation =
global_list['RMS_expectation']
        rho = parameter_list['rho']
        delta = parameter_list['delta']
        gradient = compute_gradient(x, y, w)
        expectation = rho * expectation + (1 - rho)
* (gradient ** 2)
        global_list['RMS_expectation'] =
expectation

        w -= lr * gradient / (np.sqrt(expectation)
+ delta)
    if method == 'ADAM':

```

```

        # adam_para_list = {'delta':10e-8,
'rho1':0.9, 'rho2':0.999, 'learning_rate':0.1}

        delta = parameter_list['delta']
        rho1 = parameter_list['rho1']
        rho2 = parameter_list['rho2']
        lr = parameter_list['learning_rate']
        s = global_list['ADAM_s']
        r = global_list['ADAM_r']
        gradient = compute_gradient(x, y, w)
        s = rho1 * s + (1 - rho1) * gradient
        r = rho2 * r + (1 - rho2) * (gradient **
2)

        s_hat = s / (1 - rho1)
        r_hat = r / (1 - rho2)
        w -= (lr * s_hat) / (np.sqrt(r_hat) + delta)
        global_list['ADAM_s'] = s
        global_list['ADAM_r'] = r
    if method == 'Adadelata':
        # adadelata_para_list = {'delta':10e-7,
'miu':0.9, 'learning_rate':0.1}

        r = global_list['ADADELDA_r']
        e = global_list['ADADELDA_e']

```



```

        miu = parameter_list['miu']
        delta = parameter_list['delta']
        lr = parameter_list['learning_rate']
        grad = compute_gradient(x, y, w)
        r = miu * r + (1 - miu) * (grad ** 2)
        w -= (lr * grad * np.sqrt(e + delta)) /
(np.sqrt(r + delta))

        e = miu * e + (1 - miu) * (w ** 2)
        global_list['ADADELDA_r'] = r
        global_list['ADADELDA_e'] = e

if __name__ == '__main__':
    x_train, y_train =
ds.load_svmlight_file('./data/a9a')
    x_test, y_test =
ds.load_svmlight_file('./data/a9a.t')

    num_samples, num_features = x_train.shape
    num_test_samples, num_test_features =
x_test.shape

```

```
x_train = x_train.toarray()
temp = np.ones(shape=[num_samples, 1],
dtype=np.float32)
x_train = np.concatenate([x_train, temp],
axis=1)
x_test = x_test.toarray()
temp = np.zeros(shape=[num_test_samples, 1],
dtype=np.float32)
temp1 = np.ones(shape=[num_test_samples, 1],
dtype=np.float32)
x_test = np.concatenate([x_test, temp, temp1],
axis=1)
y_train = y_train.reshape([len(y_train), 1])
y_test = y_test.reshape([len(y_test), 1])

def shuffle_train():
    global x_train, y_train
    rng_state = np.random.get_state()
    np.random.shuffle(x_train)
    np.random.set_state(rng_state)
    np.random.shuffle(y_train)
```

```

batch_size = 1024
data_size = num_samples

def feed_data(batch_count):
    if (1 + batch_count) * batch_size <=
data_size:
        feed_dict = {'x': x_train[batch_count
* batch_size:(batch_count + 1) * batch_size],
                    'y': y_train[batch_count *
batch_size:(batch_count + 1) * batch_size]}
    else:
        feed_dict = {'x': x_train[batch_count
* batch_size:data_size],
                    'y': y_train[batch_count *
batch_size:data_size]}
    return feed_dict

methods = ['SGD', 'NAG', 'RMSProp',
'Adadelta', 'ADAM']
sgd_para_list = {'learning_rate': 0.01}
nag_para_list = {'miu': 0.9, 'learning_rate':

```

```

0.01}

    rmsprop_para_list = {'delta': 10e-7, 'rho':
0.95, 'learning_rate': 0.1}

    adam_para_list = {'delta': 10e-8, 'rho1': 0.9,
'rho2': 0.999, 'learning_rate': 0.1}

    adadelta_para_list = {'delta': 10e-7, 'miu':
0.9, 'learning_rate': 0.1}

    para_list = {'SGD': sgd_para_list, 'NAG':
nag_para_list, 'RMSProp': rmsprop_para_list,
                'Adadelta': adadelta_para_list,
'ADAM': adam_para_list}

    for method in methods:

        iters = []
        test_errors = []
        w = np.random.rand(feature_size+1, 1)
        parameter_list = para_list[method]
        count = 0

        for i in range(0, 3):

            shuffle_train()

            for batch_count in range(0,
int(data_size / batch_size) + 1):

                feed_dict = feed_data(batch_count)

```

```

        iters.append(count)
        count += 1
        optimizer(method=method,
parameter_list=parameter_list,
x=feed_dict['x'], y=feed_dict['y'], w=w)

test_errors.append(compute_loss(x_test, y_test,
w))

        plt.plot(iters, test_errors,
label=method)

plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

(Fill in the contents of 8-11 respectively for logistic regression and linear classification)

8. The initialization method of model parameters:

Randomly initializing.

9. The selected loss function and its derivatives:

$$\mathbf{w} = (\mathbf{w}; 1)$$

$$p_i = \text{sigmoid}(w^T x_i)$$

$$\text{Loss} = -\frac{1}{N} \sum_{i=1}^N y_i \log(p_i) + (1 - y_i) \log(1 - p_i)$$

$$\frac{\partial \text{Loss}}{\partial w} = -\frac{1}{N} \sum_{i=1}^N x_i (p_i - y_i)$$

10. Experimental results and curve:(Fill in this content for various methods of gradient descent respectively)

Hyper-parameter selection:

SGD:

Learning rate = 0.01

NAG:

Learning rate = 0.01

Miu = 0.9

Momentum_ini = [0,0 ...,0]^T

RMSProp:

Learning rate = 0.1

Expectation = 1

ρ = 0.95

delta=10e − 7

Adam:

delta = 10e-8

rho1 = 0.9

rho2 = 0.999

lr = 0.1

s = 0

r = 0

Adadelata

r = 0

e = 0


miu = 0.9


delta = 10e-7

lr = 10

Predicted Results (Best Results):

Loss curve:





11. Results analysis:

12. Similarities and differences between logistic regression and linear classification :

13. Summary: